Model differences, with uncertainty

Original test results | six comparisons | the selected model stays fixed

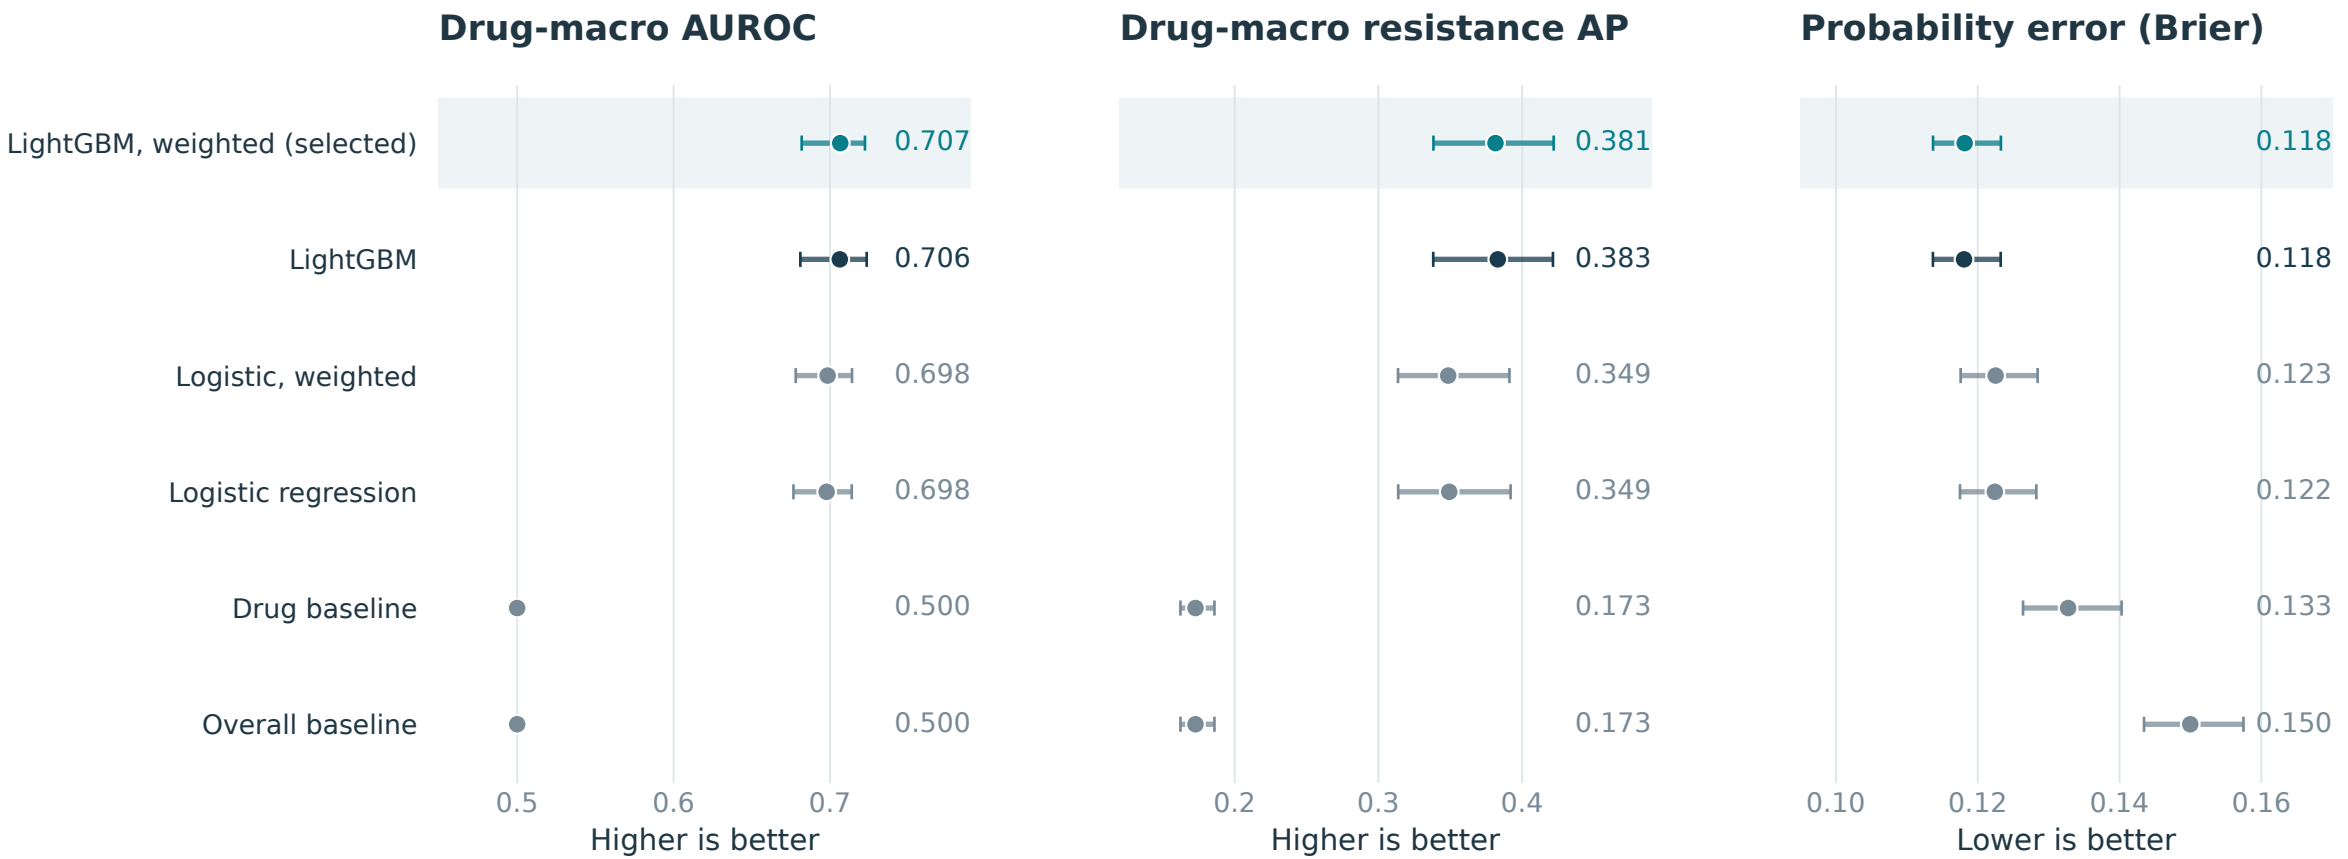

Dots show estimates; lines show 95% patient-bootstrap intervals from the primary run (200 draws). AP is average precision for resistance. Weighted and unweighted LightGBM are very close; this is not evidence of a clear benefit from class weighting.